

CSE6010 Mini-Project

Jipeng Wu

College of Computing
Georgia Institute of Technology

Project Presentation, 2015

Outline

- 1 Introduction
- 2 Problem Specification
- 3 Simulation of RTOS Runtime Context
- 4 Analysis

2 Concepts

- 1 Real-time Operating System
- 2 Task Scheduling

The Multi-level Feedback Queue

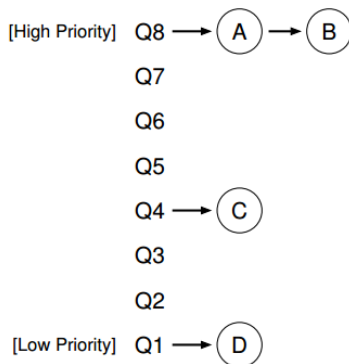


Figure: The Multi-level Feedback Queue

We Need Simple Algorithms

- 1 Sometimes RTOS must guarantee events in several nanoseconds.
- 2 Even an AND gate needs about 500 picosecond! An instruction usually needs 5ns!
- 3 INTEGRITY-178B uses RMA(Rate Monotonic Analysis) to reduce syscall cost.

We Need Simple Algorithms

- 1 Sometimes RTOS must guarantee events in several nanoseconds.
- 2 Even an AND gate needs about 500 picosecond! An instruction usually needs 5ns!
- 3 INTEGRITY-178B uses RMA(Rate Monotonic Analysis) to reduce syscall cost.

We Need Simple Algorithms

- 1 Sometimes RTOS must guarantee events in several nanoseconds.
- 2 Even an AND gate needs about 500 picosecond! An instruction usually needs 5ns!
- 3 INTEGRITY-178B uses RMA(Rate Monotonic Analysis) to reduce syscall cost.

We Need Simple Algorithms

- 1 Sometimes RTOS must guarantee events in several nanoseconds.
- 2 Even an AND gate needs about 500 picosecond! An instruction usually needs 5ns!
- 3 INTEGRITY-178B uses RMA(Rate Monotonic Analysis) to reduce syscall cost.

Periodic Hard-real-time Job Sets

- In this simulation, we have three or four periodic hard-real-time job sets:

$S_i = J_{i1}, J_{i2}, \dots, J_{in}$ Each job in S_i shares the same period p and the same instruction sequence length l .

- J_{ik} arrives at the $p(k-1)$ -th unit of time and must be finished before the pk -th unit of time.
- We need to simulate different runs under different real-time schedulers and evaluate their performance.

Periodic Hard-real-time Job Sets

- In this simulation, we have three or four periodic hard-real-time job sets:
 $S_i = J_{i1}, J_{i2}, \dots, J_{in}$ Each job in S_i shares the same period p and the same instruction sequence length l .
 - J_{ik} arrives at the $p(k-1)$ -th unit of time and must be finished before the pk -th unit of time.
- We need to simulate different runs under different real-time schedulers and evaluate their performance.

Periodic Hard-real-time Job Sets

- In this simulation, we have three or four periodic hard-real-time job sets:
 $S_i = J_{i1}, J_{i2}, \dots, J_{in}$ Each job in S_i shares the same period p and the same instruction sequence length l .
 - J_{ik} arrives at the $p(k-1)$ -th unit of time and must be finished before the pk -th unit of time.
- We need to simulate different runs under different real-time schedulers and evaluate their performance.

Periodic Hard-real-time Job Sets

- In this simulation, we have three or four periodic hard-real-time job sets:
 $S_i = J_{i1}, J_{i2}, \dots, J_{in}$ Each job in S_i shares the same period p and the same instruction sequence length l .
 - J_{ik} arrives at the $p(k-1)$ -th unit of time and must be finished before the pk -th unit of time.
- We need to simulate different runs under different real-time schedulers and evaluate their performance.

Periodic Hard-real-time Job Sets

- In this simulation, we have three or four periodic hard-real-time job sets:
 $S_i = J_{i1}, J_{i2}, \dots, J_{in}$ Each job in S_i shares the same period p and the same instruction sequence length l .
 - J_{ik} arrives at the $p(k-1)$ -th unit of time and must be finished before the pk -th unit of time.
- We need to simulate different runs under different real-time schedulers and evaluate their performance.

RTOS Schedulers And Periodic Jobs

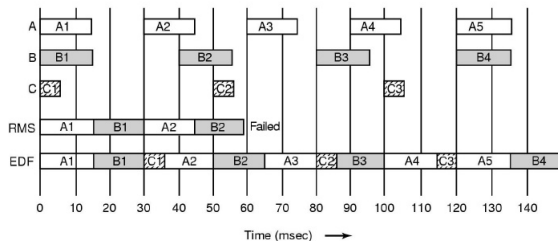


Figure: EDF and RMS

RT OS Simulation

- 1 Define tasks as sequences of instructions and a pointer to where it is executing. Each instruction costs the same constant units of time.
- 2 Use a paralleled thread to maintain `time_tick`, continuously incrementing it. Schedulers can read it.
- 3 Simulate periodic job events: every `p` units of time, add a new job into the task list, and record the execution time of last job.
- 4 For simplicity, there is only one interruption in our RTOS, the clock interruption.
- 5 The architecture allows both preemptive and non-preemptive scheduling.
- 6 Global resources: tick, task list, current task pointer, PC.

RT OS Simulation

- 1 Define tasks as sequences of instructions and a pointer to where it is executing. Each instruction costs the same constant units of time.
- 2 Use a paralleled thread to maintain `time_tick`, continuously incrementing it. Schedulers can read it.
- 3 Simulate periodic job events: every `p` units of time, add a new job into the task list, and record the execution time of last job.
- 4 For simplicity, there is only one interruption in our RTOS, the clock interruption.
- 5 The architecture allows both preemptive and non-preemptive scheduling.
- 6 Global resources: tick, task list, current task pointer, PC.

RT OS Simulation

- 1 Define tasks as sequences of instructions and a pointer to where it is executing. Each instruction costs the same constant units of time.
- 2 Use a paralleled thread to maintain `time_tick`, continuously incrementing it. Schedulers can read it.
- 3 Simulate periodic job events: every `p` units of time, add a new job into the task list, and record the execution time of last job.
- 4 For simplicity, there is only one interruption in our RTOS, the clock interruption.
- 5 The architecture allows both preemptive and non-preemptive scheduling.
- 6 Global resources: tick, task list, current task pointer, PC.

RT OS Simulation

- 1 Define tasks as sequences of instructions and a pointer to where it is executing. Each instruction costs the same constant units of time.
- 2 Use a paralleled thread to maintain `time_tick`, continuously incrementing it. Schedulers can read it.
- 3 Simulate periodic job events: every `p` units of time, add a new job into the task list, and record the execution time of last job.
- 4 For simplicity, there is only one interruption in our RTOS, the clock interruption.
- 5 The architecture allows both preemptive and non-preemptive scheduling.
- 6 Global resources: tick, task list, current task pointer, PC.

RT OS Simulation

- 1 Define tasks as sequences of instructions and a pointer to where it is executing. Each instruction costs the same constant units of time.
- 2 Use a paralleled thread to maintain `time_tick`, continuously incrementing it. Schedulers can read it.
- 3 Simulate periodic job events: every `p` units of time, add a new job into the task list, and record the execution time of last job.
- 4 For simplicity, there is only one interruption in our RTOS, the clock interruption.
- 5 The architecture allows both preemptive and non-preemptive scheduling.
- 6 Global resources: `tick`, task list, current task pointer, `PC`.

RT OS Simulation

- 1 Define tasks as sequences of instructions and a pointer to where it is executing. Each instruction costs the same constant units of time.
- 2 Use a paralleled thread to maintain `time_tick`, continuously incrementing it. Schedulers can read it.
- 3 Simulate periodic job events: every `p` units of time, add a new job into the task list, and record the execution time of last job.
- 4 For simplicity, there is only one interruption in our RTOS, the clock interruption.
- 5 The architecture allows both preemptive and non-preemptive scheduling.
- 6 Global resources: tick, task list, current task pointer, PC.

Clock Interrupt Simulation Thread

Listing 1: clock interruption ISR simulation

```
while(true){ //ISR(Interruption Service Routine)
    tick++;
    delay(CLOCK_CYCLE); //long enough to print
                        something
}
```

Kernel And User Code Simulation Thread

Listing 2: main simulation thread

```
while(true){  
    if(interrupted){ //simulate kernel-mode code  
        store_context_of(current_task); //store  
            context  
        if(is_preemptive){  
            schedule();  
        }  
        else{  
            time_update();  
        }  
    }  
    else{ // runtime strictly < 1 clock cycle  
        exec(PC); //spurious exec function, print sth.  
    }  
}
```

Tick Updation Routine

Listing 3: main simulation thread

```
void timeupdate() {  
    current_task->inst[PC].tick--;  
    if (!current_task->inst[PC].tick) {  
        PC++;  
    }  
    if (PC==current_task->nr_inst) {  
        schedule();  
    }  
}
```

Performance Analysis

The analysis will include:

- 1 The most important performance metric is the Successive Ratio and it is defined as,

$$SR = \frac{\text{Number of jobs successfully scheduled}}{\text{Total number of jobs arrived}}$$

- 2 Execution Time E_i .

$$E_i = \begin{cases} \text{execution time of the job} & \text{if deadline met} \\ 0 & \text{if deadline not met} \end{cases}$$

- 3 ECU(Effective CPU Utilization) gives information about how efficiently the processor is used and it is defined as,

$$ECU = \sum_{i \in R} \frac{E_i}{T}. \quad T = \text{Total time of scheduling}$$

Performance Analysis

The analysis will include:

- 1 The most important performance metric is the Successive Ratio and it is defined as,

$$SR = \frac{\text{Number of jobs successfully scheduled}}{\text{Total number of jobs arrived}}$$

- 2 Execution Time E_i .

$$E_i = \begin{cases} \text{execution time of the job} & \text{if deadline met} \\ 0 & \text{if deadline not met} \end{cases}$$

- 3 ECU(Effective CPU Utilization) gives information about how efficiently the processor is used and it is defined as,

$$ECU = \sum_{i \in R} \frac{E_i}{T}. \quad T = \text{Total time of scheduling}$$

Performance Analysis

The analysis will include:

- 1 The most important performance metric is the Successive Ratio and it is defined as,

$$SR = \frac{\text{Number of jobs successfully scheduled}}{\text{Total number of jobs arrived}}$$

- 2 Execution Time E_i .

$$E_i = \begin{cases} \text{execution time of the job} & \text{if deadline met} \\ 0 & \text{if deadline not met} \end{cases}$$

- 3 ECU(Effective CPU Utilization) gives information about how efficiently the processor is used and it is defined as,

$$ECU = \sum_{i \in R} \frac{E_i}{T}. \quad T = \text{Total time of scheduling}$$

Performance Analysis

The analysis will include:

- 1 The most important performance metric is the Successive Ratio and it is defined as,

$$SR = \frac{\text{Number of jobs successfully scheduled}}{\text{Total number of jobs arrived}}$$

- 2 Execution Time E_i .

$$E_i = \begin{cases} \text{execution time of the job} & \text{if deadline met} \\ 0 & \text{if deadline not met} \end{cases}$$

- 3 ECU(Effective CPU Utilization) gives information about how efficiently the processor is used and it is defined as,

$$ECU = \sum_{i \in R} \frac{E_i}{T}. \quad T = \text{Total time of scheduling}$$

Q & A